

# **Explanation Report**

## **AI Quant Portfolio: LSTM vs Transformer**

Prepared by: Shivam Singh (project)  
Period: 2010-01-01 to 2020-12-31

Deliverables included:

- Code repository (src/, main.py)
- README and run instructions
- Results and figures (results/)
- This explanation report (PDF)

# Approach & Methodology

## Approach & Methodology

### 1. Data Collection & Preprocessing

- Historical daily prices were downloaded via Yahoo Finance (yfinance).
- Selected assets: global indices / ETFs (S&P500, FTSE, Nikkei, EEM), Gold, and U.S. 10Y proxy.
- Cleaned and aligned time series; computed log-returns and rolling features (MA(5), MA(21), vol21).

### 2. Modeling

- Trained two models to forecast next-day returns per asset:
  - a) LSTM (2 layers, 50 hidden units, dropout=0.2, lookback=60 days)
  - b) Transformer (2 encoder layers, 8 heads, model\_dim=64, lookback=60 days)
- Training objective: MSE on next-day returns. Early stopping used on validation loss.

### 3. Portfolio Construction

- Each day, predicted next-day returns were input into a Mean-Variance optimizer (max Sharpe / tangency portfolio, long-only,  $0 \leq w \leq 1$ ,  $\sum(w)=1$ ).
- Covariance estimated using a rolling 60-day window of historical returns.
- Portfolios rebalanced daily; turnover cost set as 0.1% per unit traded.

### 4. Backtest & Evaluation

- Simulated daily NAV for each strategy and computed metrics: annualized return, annualized volatility, Sharpe ratio, and maximum drawdown.

## Description of Models Used

### Description of Models Used

#### LSTM (Long Short-Term Memory):

- A recurrent neural network architecture designed to capture sequential dependencies.
- Configured with 2 stacked LSTM layers with 50 hidden units each, using the last timestep output followed by a dense layer to predict multi-asset returns.

#### Transformer (Encoder-only):

- Uses self-attention to capture relationships across time steps without recurrent connections.
- Configured as an encoder stack with 2 layers, 8 attention heads, model dimension 64 and FFN=128.
- Input sequence is projected to model\_dim, passed through transformer encoders, and the final token (last timestep) is used for predictions.

#### Why both:

- LSTM is a classical time-series model good at local sequential patterns; Transformer captures global interactions and often generalizes better on longer dependencies.

## Challenges Faced & How They Were Addressed

### Challenges Faced & How They Were Addressed

#### 1. Data compatibility with yfinance:

- Issue: yfinance changed defaults and sometimes 'Adj Close' wasn't present.
- Fix: used a robust loader that falls back to 'Close' column and handles missing tickers.

#### 2. Indexing / slicing bugs during sequence creation:

- Issue: using `df[...]` semantics attempted column access instead of row slicing.
- Fix: switched to explicit positional indexing with `.iloc` for creating sequences.

#### 3. Numerical issues (log of zero or inf):

- Issue: `log(0)` appeared when `pct_change` produced zeros, leading to `-inf` / `NaN`.
- Fix: used `np.log1p` on returns and sanitized infinities and NaNs before covariance estimation.

#### 4. Module import paths when running from project root:

- Issue: Python couldn't find modules when running `main.py`.
- Fix: made `src` a package (added `__init__.py`) and used relative imports.

#### 5. Dependency problems (empyrical):

- Issue: `empyrical` failed to install on newer Python versions.
- Fix: removed `empyrical` from requirements and implemented required metrics in `utils.py`.

Backtest Results (2010-01-01 to 2020-12-31)

	Transformer	LSTM
Annual Return	15.74%	0.55%
Annual Volatility	22.42%	23.35%
Sharpe Ratio	0.702	0.023
Max Drawdown	-0.425	-0.490

# Transformer vs LSTM: Annual Return & Sharpe

